

Comparative Analysis of AI-Assisted and Manual Data Labeling Accuracy Across Scientific Datasets

Abhay Kumar Awasthi

Introduction

The unprecedented proliferation of high-dimensional data across scientific disciplines has established data annotation as the primary bottleneck in the development, validation, and deployment of supervised machine learning systems. For decades, the foundational bedrock of artificial intelligence has been the "ground truth" provided by manual human annotation. In highly specialized domains such as computational pathology, single-cell transcriptomics, clinical trial management, geospatial remote sensing, and astrophysics, this manual process demands significant domain expertise. Consequently, it is prohibitively expensive, inherently unscalable, and notoriously susceptible to human fatigue, cognitive bias, and inter-observer variability. The historical assumption that expert manual labeling represents an infallible gold standard has been increasingly challenged by empirical evidence demonstrating that even highly curated benchmark datasets contain substantial rates of label noise and systematic bias.

To address these critical inefficiencies, the scientific community is rapidly shifting toward AI-assisted data labeling. This paradigm operates on a Human-in-the-Loop (HITL) framework, wherein machine learning algorithms—ranging from foundation models and large language models (LLMs) to active learning heuristics and zero-shot classifiers—provide initial annotations, segmentations, or localized bounding boxes.




Figure 1: Conceptual Architecture of the Human-in-the-Loop AI-Assisted Scientific Data Labeling Pipeline.

As illustrated in Figure 1, these algorithmic suggestions are subsequently reviewed, refined, or adjudicated by human annotators. This comprehensive comparative analysis examines the accuracy, efficiency, statistical reliability, and cognitive implications of AI-assisted versus manual data labeling across diverse scientific datasets. By synthesizing empirical evidence from medical imaging, genomics, clinical text processing, and astronomy, this report elucidates the transformative potential of hybrid intelligence while critically evaluating the emergent risks of automation bias, out-of-distribution shifts, and the propagation of label noise.

1 Theoretical and Statistical Frameworks for Ground Truth

Before directly comparing manual and AI-assisted modalities, it is necessary to establish the mathematical and statistical frameworks used to define "accuracy" in contexts where absolute objective truth is fundamentally inaccessible. In scientific datasets, ground truth is typically established through consensus among multiple domain experts. The reliability of this consensus is quantified using Inter-Annotator Agreement (IAA) metrics, which evaluate the consistency with which multiple annotators assign categories, segment boundaries, or parse textual relationships in a shared dataset.

1.1 Rethinking Annotation as a Measurement Process

Machine learning research has conventionally treated all annotator disagreement as noise. However, the Human Label Variation (HLV) framework posits that annotation is a highly complex measurement process where label variation stems from four distinct sources: instance difficulty, annotator bias, situational noise, and relational alignment. Decomposing these outcomes reveals that what algorithms perceive as "noise" is often valid, subjective medical or scientific disagreement over ambiguous phenomena. When researchers blindly enforce consensus (e.g., via simple majority voting) on ambiguous items, they obscure these critical distinctions and limit the downstream model's ability to navigate edge cases. A study involving eleven critical care consultants annotating Intensive Care Unit (ICU) data revealed only "fair" agreement (Fleiss' $\kappa = 0.383$) internally, and "minimal" agreement ($\kappa = 0.255$) when externally validated against out-of-network datasets, demonstrating that a singular "super expert" rarely exists in acute clinical settings.

1.2 Quantifying Inter-Annotator Agreement

To rigorously evaluate annotation quality, researchers rely on metrics that correct for chance agreement, distinguishing between true scientific consensus and random overlap. Cohen's Kappa (κ) is the standard metric for binary or categorical classification tasks involving two annotators. It is formulated mathematically as:

$$\kappa = \frac{P_o - P_e}{1 - P_e} \quad (1)$$

Where P_o represents the relative observed agreement among raters, and P_e represents the hypothetical probability of chance agreement. In complex medical annotation tasks, such as classifying oral potentially malignant disorders from intraoral images, manual expert agreement often reaches $\kappa = 0.90$ indicating excellent reliability.

For datasets requiring annotation by more than two experts, Fleiss' Kappa extends the mathematical foundation of Cohen's Kappa to assess the consistency of multiple raters evaluating the same items simultaneously. Krippendorff's Alpha (α) provides an even more robust and flexible formulation, accommodating nominal, ordinal, interval, and ratio data, while successfully handling missing values—a common occurrence in highly sparse scientific datasets.

In computer vision tasks, particularly medical image segmentation and satellite remote sensing, accuracy cannot be evaluated using categorical agreement alone; it requires spatial

overlap metrics. The Dice Similarity Coefficient (DSC) and Intersection over Union (IoU) serve as the primary standards:

$$\text{DSC} = \frac{2|X \cap Y|}{|X| + |Y|} \quad (2)$$

$$\text{IoU} = \frac{|X \cap Y|}{|X \cup Y|} \quad (3)$$

Where X and Y represent the spatial masks generated by different human annotators or by an AI model against a human reference. To aggregate varying human annotations into a single, unified ground truth mask for AI training, algorithms such as Simultaneous Truth and Performance Level Estimation (STAPLE) are deployed. STAPLE computes a probabilistic estimate of the true segmentation based on the historical performance level and bias of each individual annotator, yielding a highly resilient training target.

Table 1: Inter-Annotator Agreement (IAA) Metrics

IAA Metric	Mathematical Characteristics	Primary Application	Applications	Strengths in Scientific Labeling
Cohen’s Kappa	Binary/Categorical	Corrects for chance agreement		Standard for dual-annotator verification (e.g., textual triage).
Fleiss’ Kappa	Multi-rater Categorical	Generalization of Cohen’s for n raters		Ideal for crowd-sourced or multi-expert panel consensus.
Krippendorff’s Alpha	Complex/Mixed Data	Handles missing data and various scales		Highly robust for sparse datasets with incomplete overlapping reviews.
STAPLE	Spatial Segmentation	Expectation-Maximization algorithm	algorithm	Aggregates multiple bounding boxes/masks into a true probabilistic label.
Dice / IoU	Spatial Overlap	Harmonic mean of precision and recall		Standard for evaluating AI boundaries against human baseline masks.

2 Medical Imaging and Computational Pathology

The integration of artificial intelligence into medical image annotation has demonstrated profound improvements in both processing speed and boundary precision. Medical imaging requires pixel-perfect accuracy, as downstream clinical models utilize these labels for oncology diagnostics, surgical planning, and anomaly detection.

2.1 Foundation Models in Radiology: SAM and MedSAM

The Segment Anything Model (SAM), a foundation model trained on over one billion mask annotations in the natural image domain, catalyzed a shift from specialist segmentation net-

works to prompt-driven generalist models. However, the direct application of zero-shot SAM to medical imaging often yields suboptimal accuracy due to the vast domain gap between natural photography and complex radiological scans.

To bridge this gap, domain-adapted models such as MedSAM have been fine-tuned on extensive medical datasets encompassing 1,570,263 image-mask pairs across 10 imaging modalities and over 30 cancer types. Comparative studies evaluating zero-shot inference for lumbar spine MRI segmentation reveal a distinct hierarchy in annotation accuracy. When segmenting intervertebral discs (IVDs) and vertebral bodies (VBs), standard SAM achieved DSCs of 0.64 and 0.83, respectively. MedSAM significantly outperformed the baseline with DSCs of 0.79 and 0.88. Despite this marked improvement, fully supervised, task-specific architectures like nnU-Net continue to represent the upper boundary of clinical accuracy, achieving a DSC of 0.99 for the same anatomical structures. Similarly, evaluations on the BraTS 2023 adult glioma and pediatrics dataset indicated that while SAM and SAM 2 can achieve high Dice scores (up to 0.894) when provided with perfectly precise bounding box prompts, nnU-Net remains dominant due to the practical impossibility of humans consistently providing such perfect prompts during routine clinical workflows.

The architectural advantage of AI-assisted foundation models does not necessarily lie in surpassing the accuracy of specialized models like nnU-Net, but in their capacity to accelerate the generation of training data for new tasks without requiring prior labeled datasets. SAM’s architecture consists of an image encoder, a prompt encoder, and a mask decoder, allowing for extreme flexibility in how humans interact with the image. Recent advancements have introduced the Box Predictor module into the MedSAM architecture. This lightweight module estimates an approximate bounding box from a single user click using localized image embedding features, providing spatial guidance that reduces the ambiguity of point prompts while introducing only 1.6M additional parameters. Across diverse datasets (FLARE22, BRISC, BUSI, LungSegDB), this approach achieved impressive Dice scores ranging from 0.88 to 0.98.

Furthermore, utilizing Granular Box Prompts (GB-SAM)—where clinicians draw localized bounding boxes around regions of interest rather than outlining complex polygons—drastically minimizes human effort while raising accuracy. In histopathology datasets (e.g., CRAG), GB-SAM achieved a DSC of 0.885 utilizing only 25% of the training data, effectively outperforming traditional U-Net architectures (DSC 0.857) trained on identical subsets.

2.2 Computational Pathology and Annotation Throughput

In computational pathology, Whole Slide Images (WSI) contain billions of pixels, making manual annotation excessively labor-intensive and prone to fatigue. A benchmark study analyzing different annotation peripherals and methodologies demonstrated the extreme efficiency gains of AI-assisted workflows. Pure manual annotation using a standard computer mouse averaged 29.9 minutes per slide with a high 24% inter-observer variability. Touchpad annotation was the least efficient, averaging 47.5 minutes with a 45% variability. Conversely, a semi-automated, AI-assisted approach combining initial manual interactions with SAM reduced the average annotation time to 13.6 minutes while simultaneously achieving the lowest inter-observer variability at just 2%.

Similarly, AI-collaborative labeling (AICL) frameworks deployed for complex volumetric tasks, such as CT intracavitary hemorrhage segmentation, demonstrate substantial geometric improvements. AICL reduced the manual time effort by 2.8-fold compared to standard clinical staff editing and 8.7-fold compared to expert radiologist labeling from scratch. The qualitative Likert grades awarded to the final annotations were significantly higher for AI-assisted labels (8.4) compared to purely manual expert labels (7.8), suggesting that algorithmic assistance uncovers boundary details that human fatigue naturally obscures over long sessions.

2.3 Hybrid Intelligence and Uncertainty-Driven Workflows

To optimize the balance between automated efficiency and human clinical judgment, "HybridMS" (Hybrid Medical Segmentation) frameworks have been introduced. These systems rely on uncertainty-driven feedback mechanisms that evaluate the AI's confidence in its own segmentation. If the model determines a region is statistically ambiguous (e.g., poorly contrasted tumor boundaries), it selectively triggers a human clinician for input. In a lung segmentation trial for tuberculosis detection on chest X-rays, HybridMS achieved a Dice score of 0.9538 (compared to a baseline MedSAM of 0.9435), while reducing standard annotation time by 82% and challenging case annotation time by 60%. This selective intervention maximizes throughput without compromising regulatory or clinical safety standards.

Table 2: Performance and Efficiency Across Modalities

Modality / Architecture	Clinical Application	Performance Metric (Dice)	Efficiency Gains over Manual Baseline
Zero-Shot SAM	Lumbar MRI	DSC: 0.64-0.83	High (Requires exact prompts)
MedSAM	Lumbar MRI	DSC: 0.79-0.88	High (Lower prompt dependency)
nnU-Net (Supervised)	Lumbar/Glioma	DSC: 0.99	None (Requires massive pre-labeled data)
GB-SAM (Box Prompt)	Histopathology (CRAG)	DSC: 0.885	Achieved with only 25% of required data
Semi-Automated SAM	WSI Pathology	2% Variability	Time reduced from 29.9 min to 13.6 min
HybridMS (Uncertainty)	Chest X-ray (Tuberculosis)	DSC: 0.9538	Annotation time reduced by 82%

3 Genomics, Single-Cell Transcriptomics, and Clinical Text Processing

Beyond visual segmentation, a vast proportion of scientific datasets consist of unstructured clinical text, concomitant medication records, adverse event narratives, and high-dimensional genomic sequences.

3.1 Single-Cell RNA Sequencing (scRNA-seq) Annotation

Single-cell RNA sequencing (scRNA-seq) provides unprecedented resolution into cellular heterogeneity, allowing researchers to build highly detailed transcriptomic atlases. Traditionally, cell type annotation has relied heavily on manual intervention. Domain experts compare differentially expressed genes in computationally generated clusters against canonical marker genes documented in literature. This manual approach is highly subjective, labor-intensive, and

lacks standardized criteria across different laboratories, leading to fragmented and incompatible datasets.

The introduction of specialized Large Language Models and multi-agent AI systems has fundamentally restructured bioinformatics pipelines. Advanced benchmarking of leading LLMs (e.g., Kimi-k2, GPT-5, Claude-4.1, Grok-4) across 34 diverse human and mouse datasets demonstrated that AI models profoundly outperform traditional computational methods (such as SingleR, Cellmarker2.0, and ScType), particularly in their ability to resolve fine-grained, rare cell subtypes. These capabilities are integrated into platforms like DeepCellSeek, which leverage statistical significance and log-fold change rankings to optimize the prompts fed into LLMs.

The "Tabula Sapiens" dataset, a massive multi-organ transcriptomic atlas, serves as a primary benchmark for evaluating AI annotation accuracy against expert manual annotations. Utilizing packages like AnnDictionary (a framework integrating LangChain and Scanpy), researchers found that LLM-driven *de novo* cell-type annotation outperformed earlier computational iterations by approximately 20% (achieving 77% accuracy compared to 58%). The analysis indicates that for major cell types, modern LLMs routinely exceed 80-90% agreement with manual expert ground truth.

Further advancements involve agentic systems like CASSIA, which leverage Claude-3.5 to achieve superior performance metrics. In comprehensive tests across 970 cell types from classical benchmark datasets (including GTEX, Tabula Sapiens, Human Cell Landscape, and Mouse Cell Atlas), CASSIA achieved an exceptional accuracy score of 0.92. AI-assisted genomic annotation platforms not only deliver labels but increasingly provide explainable reasoning—detailing the exact marker genes and literature citations used to classify a cluster—thereby guarding against algorithmic hallucination and allowing human bioinformaticians to confidently validate the model’s logic. Tools like CytoAnalyst follow a similar paradigm, blending unsupervised Louvain clustering with AI-assisted inference, but maintaining that ultimate accuracy requires the researcher to act as an adjudicator in the loop.

3.2 Clinical Trial Data Cleaning and Extraction

In pharmaceutical drug development, cleaning clinical trial data represents a severe operational bottleneck. The manual reconciliation of these datasets is highly susceptible to human error due to repetitive task fatigue and the inability to scale with exponentially growing data volumes.

In a controlled experimental study involving experienced medical reviewers, the implementation of "Octozi"—an AI-assisted platform combining Large Language Models with domain-specific clinical heuristics—yielded unprecedented improvements. The AI-assisted workflow increased data cleaning throughput by 6.03-fold. More importantly, the system caused cleaning errors to plummet from a manual baseline of 54.67% to just 8.48%, representing a 6.44-fold improvement in accuracy.

A persistent issue in manual clinical trial reviews is the generation of "false positive queries," where reviewers flag data anomalies that are not clinically meaningful, placing an unnecessary operational burden on trial sites. The AI-assisted system reduced false positive queries by a staggering 15.48-fold. Economically, projecting these efficiency and accuracy gains onto a representative Phase III oncology trial revealed potential cost savings of \$5.1 million, driven largely by accelerating the database lock timeline by 33% (a 5-day reduction saving \$4.4 million) and reducing query management burdens.

Similarly, in evidence synthesis and pharmacovigilance, artificial intelligence vastly outperforms double-human data extraction. Randomized controlled trials (RCTs) evaluating data extraction strategies for systematic reviews in sleep medicine indicate that hybrid approaches—where AI performs the initial extraction followed by human verification—drastically reduce the hours required while retaining high fidelity. In the curation of precision oncology knowledgebases like CIVIC-Fact, AI tools streamline the extraction of genetic variants and drug responses

from literature, maintaining substantial inter-annotator agreement (Fleiss’ $\kappa = 0.65$, 77.9% raw agreement) while reducing the staff hours required for primary literature review.

4 Astrophysics and Geospatial Remote Sensing

In disciplines characterized by massive planetary or cosmic scale, such as Earth observation and astrophysics, the sheer volume of high-resolution imagery necessitates automated pipelines. However, labeling this data requires nuanced spatial and morphological reasoning that challenges basic computer vision architectures.

4.1 Galaxy Morphology Classification and Green AI

The crowdsourced citizen science project Galaxy Zoo historically demonstrated the power of distributed manual annotation, where volunteers classified millions of galaxies into morphological categories (e.g., spiral, elliptical, merging) through complex decision trees. As telescope capabilities advance—such as with the Rubin Observatory Legacy Survey of Space and Time (LSST) and the Euclid space mission—the volume of data renders human crowdsourcing mathematically insufficient.

Deep learning models pre-trained on Galaxy Zoo labels, such as the Zoobot architecture, have successfully automated the detection of complex morphological phenomena. In the identification of lopsided spiral galaxies versus symmetric galaxies, fine-tuned Zoobot models achieved a testing accuracy of 87% against expert-verified ground truth, successfully labeling datasets that would take human crowds years to process. Similarly, YOLOv5 architectures have been deployed to detect edge-on galaxies, achieving detection rates approaching 97% over vast survey data like Pan-STARRS. Modern approaches are transitioning away from astronomy-specific supervised models toward general-purpose Vision-Language Models (VLMs) and Visual Question Answering (VQA) frameworks. By distilling knowledge from general VLMs, researchers can generate highly accurate morphological captions and bounding boxes without requiring massive new tranches of manually labeled astronomical data.

Crucially, the scale of astronomical AI brings environmental and computational constraints into focus. A large-scale "Green AI" benchmark for galaxy morphology classification revealed that while ensemble forests and deep learning models achieve similar accuracy, a GPU-optimized multilayer perceptron (MLP) achieved state-of-the-art accuracy of 98% while emitting 20 times less CO₂ than competing ensemble methods. This indicates that AI-assisted annotation pipelines must optimize not just for accuracy, but for inference latency and hardware-algorithm co-design to remain sustainable.

4.2 Geospatial Remote Sensing and Earth Foundation Models

In remote sensing, AI-assisted annotation involves parsing satellite imagery for land cover classification, infrastructure damage, and ecological monitoring. Earth Foundation Models (FMs) are driving a paradigm shift from traditional numerical modeling to AI-assisted predictive meteorology and spatial reasoning. Spatiotemporal forecasting models are now capable of moving beyond deterministic predictions to unified, physics-informed simulations. For instance, the Depth Any Canopy (DAC) model fine-tunes a monocular depth foundation model to estimate forest canopy height from aerial imagery, achieving superior accuracy compared to traditional remote sensing baselines.

Semantic segmentation models, enhanced by human-in-the-loop workflows, provide continuous monitoring of surface reflectance and ecosystem regeneration. Frameworks like HealingFactor and EcoLogic utilize satellite imagery to forecast post-disturbance regeneration, evaluating causal versus correlational reasoning within complex environments. AI-assisted vectorization drastically outperforms manual Geographic Information System (GIS) delineation in terms

of speed, though human oversight remains essential for resolving edge cases in multispectral anomalies (e.g., differentiating between specific types of cloud shadow, karst regions, and surface water) where Light Detection and Ranging (LiDAR) data is unavailable. In infrastructure assessment, models leveraging 3D Ground Penetrating Radar (GPR) and AI-assisted segmentation identify concrete bridge damage and asphalt defects with precision that manual evaluation cannot match at scale. Furthermore, tools like MO-SAM (Mining-Oriented Segment Anything Model) highlight the necessity of multidisciplinary human expertise to assess the accuracy of AI-assisted remote sensing, particularly for monitoring large-scale mining activities where legal and environmental compliance is required.

5 The Psychology and Ergonomics of Human-in-the-Loop Annotation

While the quantitative benefits of AI-assisted labeling are statistically undeniable across scientific domains, the insertion of AI into the human decision-making loop introduces complex cognitive and psychological phenomena. The transition from purely manual tasks to human-AI collaboration fundamentally alters the nature of the work—shifting the human from a generative “creator” of data to a critical “verifier” of algorithmic output.

5.1 Annotation Fatigue and Cognitive Load

Manual labeling over long durations induces severe annotation fatigue, leading to a demonstrable decline in data quality. Fatigue manifests as mental exhaustion stemming from high cognitive load, repetitive visual scanning, and subjective stress when dealing with ambiguous data guidelines. Empirical observations indicate that prolonged manual annotation can drive error rates as high as 30%, rendering the data unusable for safety-critical clinical applications or precise scientific tracking.

The Human Label Variation (HLV) framework supports this, illustrating that human error is rarely random; it is highly correlated with the duration of the task and the opacity of the labeling instructions. AI-assisted pre-labeling directly mitigates this fatigue by presenting the annotator with a baseline hypothesis rather than a blank screen. The cognitive effort required to adjust an existing AI bounding box or correct a pre-filled clinical text form is significantly lower than the effort required to generate those data points from scratch. This ergonomic shift extends the duration of high-quality human attention and drastically increases overall throughput.

5.2 Automation Bias and The Anchoring Effect

The central paradox of AI assistance is that as algorithms become more highly accurate, human overseers become more complacent. This phenomenon, known as automation bias or the anchoring effect, occurs when human annotators overly rely on AI-generated suggestions, subsequently failing to detect or correct algorithmic errors. The human effectively anchors their judgment to the AI’s initial prediction, known quantitatively as the “Weight of Advice” (WoA).

In a randomized experiment involving 2,784 participants engaged in a controlled annotation task, researchers discovered counterintuitive patterns in human-AI collaboration. When workflows were strictly designed to require annotators to explicitly correct flagged AI errors, human engagement actually dropped. The tendency to blindly accept incorrect algorithmic suggestions increased, demonstrating that cognitive shortcuts dominate when the task burden is perceived as high.

Furthermore, individual psychological attitudes toward AI proved to be a stronger predictor of annotation accuracy than demographic factors or professional experience. Participants who were inherently skeptical of AI detected errors more reliably and produced higher-quality

ground truth. Conversely, annotators with highly favorable attitudes toward automation exhibited dangerous overreliance, acting as a rubber stamp for the algorithm’s mistakes. This is further highlighted by the concept of the ”Turking Test,” which examines whether LLM-generated annotations successfully adhere to established guidelines compared to human reasoning, frequently showing that AI achieves 90-95% of human performance but requires skeptical oversight for the remaining edge cases.

In clinical settings, the anchoring effect is distinctly measurable. A linear mixed-effects model (LMM) analysis in computational pathology demonstrated a statistically significant positive coefficient for AI advice, indicating moderate anchoring on the system’s output. Crucially, this anchoring bias intensified severely under time pressure, suggesting that when cognitive resources are limited, clinicians default to the AI’s judgment. To counteract automation bias, modern AI pipelines must be designed with ”explainability” and ”uncertainty quantification,” forcing the human to critically engage with the data rather than passively confirming it.

Table 3: Cognitive Mechanisms and Mitigation Strategies

Cognitive Phenomenon	Mechanism	Mitigation Strategy	Impact on Annotation Quality
Annotation Fatigue	Mental exhaustion from repetitive, high-load manual tasks	AI pre-labeling, task rotation, ergonomic UI design	Increased random errors, dropped boundaries, slower throughput
Automation Bias	Complacent trust in highly accurate AI systems	Injecting synthetic errors to test vigilance, uncertainty routing	Systematic encoding of AI hallucination into ground truth
Anchoring Effect (WoA)	Judgment skewed by the initial AI hypothesis	Blind review of low-confidence predictions, removing time constraints	Inability to correct subtle boundary deviations under time pressure

6 Mitigating Label Noise: Active Learning, Confident Learning, and OOD Detection

To maximize the efficiency of the human-AI partnership and counteract cognitive biases, advanced machine learning paradigms govern the flow of data to human annotators. Three primary frameworks dictate this routing: Active Learning (minimizing the volume of data needing annotation), Confident Learning (systematically identifying and cleaning existing label errors), and Out-of-Distribution (OOD) detection (managing domain shifts).

6.1 Active Learning and Uncertainty Sampling

Active learning operates on the principle that not all data points contribute equally to model training. Rather than randomly selecting data for human review, the algorithm actively queries the human annotator to label only the most informative or uncertain instances. This targeted methodology reduces the total annotation budget—often by 99% in dense corpora—while maintaining or exceeding the accuracy of a model trained on a fully, randomly sampled dataset.

Uncertainty sampling is the most prevalent active learning heuristic. The model calculates its predictive probability for an unlabeled sample; if the prediction hovers near the decision

boundary (e.g., maximum entropy or lowest confidence margin), the sample is flagged for human intervention. Mathematical theorems surrounding active learning demonstrate that the uncertainty upper bound in traditional active learning decays at a rate of $\mathcal{O}(B^{-\alpha/d})$, proving a significantly faster convergence rate toward the optimal predictor than random sampling.




Figure 2: Process Flow of the Active Learning Data Selection and Annotation Cycle.

As shown in Figure 2, advanced iterations, such as the ActiveAED (Active Annotation Error Detection) framework, interleave error detection directly into the active learning loop. Utilizing the Area-under-the-Margin (AUM) metric, ActiveAED continually queries humans for error corrections during the prediction loop. This dynamic feedback dramatically improves the model’s ability to flag subsequent errors, leading to gains of up to six percentage points in average precision across diverse natural language processing datasets. Efficiency is often measured using the Work Saved over Sampling (WSS) metric and the Average Time to Discovery (ATD). Simulation studies reveal that active learning models can reduce the number of publications needed to screen by 63.9% to 91.7% while still discovering 95% of relevant records (WSS@95), with ATD values indicating that researchers only need to make labeling decisions on 1.4% to 11.7% of the total dataset.

6.2 Confident Learning and Label Noise Purification

Even with meticulous human-in-the-loop oversight, real-world datasets inevitably contain label noise. Confident Learning (CL) is a model-agnostic meta-learning framework explicitly designed to identify, characterize, and remove label errors from datasets before model training.

Confident Learning operates by estimating class-conditional label noise error rates through cross-validation, subsequently employing rank-based algorithms to flag suspicious annotations. In highly sensitive domains like retinal image segmentation for diabetic retinopathy, Confident Joint Analysis (CJA) utilizing DeeplabV3+ and Prune by Noise Rate (PBNR) utilizing ResNet-18 have been successfully applied alongside human refinement. By identifying misclassified pixels and routing them to a human for re-annotation, CL-based corrections yielded measurable advancements, boosting Accuracy by 4.05% and Intersection over Union (IoU) by 9.95%. This demonstrates that algorithmic error detection coupled with human review is vastly superior to relying on unverified legacy datasets.

Furthermore, standard error detection often fails to account for semantic label bias—where error rates disproportionately impact certain demographic groups or rare edge cases. Frameworks like Decoupled Confident Learning (DeCoLe) specifically target bias-aware mislabeling, ensuring that dataset purification does not inadvertently strip away rare but valid instances, which is critically important for maintaining fairness and clinical safety in medical and social computing AI.

6.3 Out-of-Distribution (OOD) Detection and Domain Shifts

A pervasive threat to both manual and AI-assisted annotation is the domain shift or Out-of-Distribution (OOD) problem. Models trained on one dataset (e.g., MRI scans from a specific manufacturer, or text from a specific demographic) often experience severe performance degradation when applied to target data with a different statistical distribution. These shifts manifest as either covariate shifts (where the value of features changes) or semantic shifts (the emergence of novel categories not seen during training).

In AI-assisted labeling, if the model encounters OOD data and fails to recognize it as such, it will generate highly confident but entirely incorrect pre-labels. If a fatigued human annotator accepts these labels due to automation bias, the error becomes permanently encoded into the ground truth. To prevent this, OOD detection algorithms—such as LogitNorm, Monte Carlo (MC) dropout, and Deep Ensembles—are integrated directly into the annotation software. These mechanisms quantify epistemic uncertainty, allowing the system to flag inputs as “anomalous” rather than attempting to label them.

Novel approaches like Adaptive Human-Assisted (AHA) OOD learning strategically request human labels specifically within a “maximum disambiguation region,” where semantic and covariate OOD data roughly equalize. By utilizing a noisy binary search algorithm, AHA maximizes the utility of a fixed labeling budget, significantly outperforming unassisted state-of-the-art methods in both OOD generalization and detection with just a few hundred human annotations. In parallel, Distribution Shift Inversion algorithms utilize diffusion models trained only on source distributions to map OOD testing samples back toward the training distribution, bypassing the need for testing distribution requirements and providing general performance gains across multiple-domain generalization datasets.

7 Conclusion

The transition from purely manual to AI-assisted data labeling marks a fundamental structural evolution in the production of scientific knowledge. In medical imaging, foundation models like MedSAM, coupled with uncertainty-driven HybridMS architectures, significantly reduce the temporal burden of pixel-wise segmentation while preserving high levels of diagnostic fidelity. In genomics and transcriptomics, large language models organized via multi-agent systems (e.g., CASSIA) are rapidly replacing subjective manual literature reviews, resolving cellular subtypes with extraordinary accuracy. In astrophysics and remote sensing, zero-shot foundation models and Vision-Language architectures allow for the continuous, high-throughput parsing of cosmic and planetary phenomena, overcoming the physical limitations of human crowdsourcing.

However, this comprehensive analysis reveals that AI assistance is not a panacea for data quality. The integration of algorithmic pre-labeling introduces acute psychological variables—specifically, automation bias, anchoring effects, and the risk of uncritical complacency among human reviewers. While AI tools successfully mitigate annotation fatigue and decrease inter-observer variability, their ultimate efficacy is strictly tethered to the cognitive engagement of the human in the loop.

The optimal scientific annotation strategy relies heavily on Active Learning and Confident Learning frameworks that purposefully measure uncertainty. By utilizing algorithmic bounds

to strategically deploy human intellect only where it provides maximal informational gain, and aggressively filtering label noise through cross-validated algorithms, researchers can protect dataset integrity. As scientific datasets continue to expand in dimensionality and volume, hybrid intelligence—the carefully calibrated, mathematically verifiable synergy of algorithmic scale and human clinical judgment—will constitute the definitive standard for ground truth generation across all domains.

Works Cited

1. From Ground Truth to Measurement: A Statistical Framework for Human Labeling - arXiv, <https://arxiv.org/html/2604.07591v1>
2. The impact of inconsistent human annotations on AI driven clinical decision making - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC9944930/>
3. Inter-Annotator Agreement (IAA) Metrics - Emergent Mind, <https://www.emergentmind.com/topics/inter-annotator-agreement-iaa-bf6a0eff-29b6-496f-804f-9d73No8b6ed1c2>
4. Counting on Consensus: Selecting the Right Inter-annotator Agreement Metric for NLP Annotation and Evaluation - arXiv, <https://arxiv.org/html/2603.06865>
5. A Smartphone-based Comprehensive Dataset of Annotated Oral Cavity Images for Enhanced Oral Disease Diagnosis - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC13125303/>
6. A Guide to Inter-rater Reliability and Annotator Agreement in AI - iMerit, <https://imerit.ai/resources/blog/human-vs-model-agreement-how-inter-rater-consistency-shapes-b>
7. Assessing Inter-Annotator Agreement for Medical Image Segmentation - PubMed - NIH, <https://pubmed.ncbi.nlm.nih.gov/37008654/>
8. Inter-Annotator Agreement: a key metric in Labeling - Innovatiana, <https://www.innovatiana.com/en/post/inter-annotator-agreement>
9. Assessing Inter-Annotator Agreement for Medical Image Segmentation - ResearchGate, https://www.researchgate.net/publication/369076293_Assessing_Inter-Annotator_Agreement_for_Medical_Image_Segmentation
10. Segment Anything in Medical Images and Videos: Benchmark and Deployment - arXiv, <https://arxiv.org/html/2408.03322v1>
11. MedSAM2: Segment Anything in 3D Medical Images and Videos - arXiv, <https://arxiv.org/html/2504.03600v1>
12. Segment anything in medical images - PMC - NIH, <https://pmc.ncbi.nlm.nih.gov/articles/PMC10803759/>
13. MedSAM-CA: A CNN-Augmented ViT with Attention-Enhanced Multi-Scale Fusion for Medical Image Segmentation - arXiv, <https://arxiv.org/html/2506.23700v1>
14. Segment Anything Model (SAM) and Medical SAM (MedSAM) for Lumbar Spine MRI - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12196577/>
15. AI-Driven MRI-based Brain Tumour Segmentation Benchmarking - arXiv, <https://arxiv.org/pdf/2506.20786>
16. Research on Medical Image Segmentation Based on SAM and Its Future Prospects - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12189367/>

17. Enhancing MedSAM with a Lightweight Box Predictor for Medical Image Segmentation, <https://arxiv.org/html/2606.04705v2>
18. Enhancing Medical Imaging Segmentation with GB-SAM: A Novel Approach to Tissue Segmentation Using Granular Box Prompts - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11240495/>
19. Improving the Annotation Process in Computational Pathology: A Pilot Study with Manual and Semi-automated Approaches on Consumer and Medical Grade Devices - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11950598/>
20. Accelerating voxelwise annotation of cross-sectional imaging through AI collaborative labeling with quality assurance and bias mitigation - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC10362988/>
21. Hybrid intelligence in medical image segmentation - PMC - NIH, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12639002/>
22. Benchmarking large language models for cell typing in single-cell RNA-Seq - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12704438/>
23. Large Language Model Consensus Substantially Improves the Cell Type Annotation Accuracy for scRNA-seq Data — bioRxiv, <https://www.biorxiv.org/content/10.1101/2025.04.10.647852v1.full-text>
24. Benchmarking large language models for cell typing in single-cell RNA-Seq — Briefings in Bioinformatics — Oxford Academic, <https://academic.oup.com/bib/article/26/6/bbaf677/8380197>
25. Benchmarking Cell Type Annotation by Large Language Models with AnnDictionary, <https://www.biorxiv.org/content/10.1101/2024.10.10.617605v1>
26. A Benchmark Study for Large Language Models on Cell Type Annotation Using Single-Cell Data - arXiv, <https://arxiv.org/html/2412.02915v1>
27. CASSIA: a multi-agent large language model for automated and interpretable cell annotation - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12796229/>
28. CASSIA: a multi-agent large language model for reference free, interpretable, and automated cell annotation of single-cell RNA-sequencing data — bioRxiv, <https://www.biorxiv.org/content/10.1101/2024.12.04.626476v2>
29. A conceptual framework for human-AI collaborative genome annotation - PMC - NIH, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12301185/>
30. arXiv:2503.23691v1 [q-bio.GN] 31 Mar 2025, <https://arxiv.org/pdf/2503.23691>
31. CytoAnalyst web platform facilitates comprehensive single cell RNA sequencing analysis, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12328742/>
32. Leveraging AI to Accelerate Medical Data Cleaning: A Comparative Study of AI-Assisted vs. Traditional Methods - arXiv, <https://arxiv.org/pdf/2508.05519>
33. Leveraging AI to Accelerate Clinical Data Cleaning: A Comparative Study of AI-Assisted vs. Traditional Methods - arXiv, <https://arxiv.org/html/2508.05519v1>
34. Leveraging AI to Accelerate Medical Data Cleaning: A Comparative Study of AI-Assisted vs. Traditional Methods - arXiv, <https://arxiv.org/html/2508.05519>

35. Comparing the accuracy of AI-assisted data extraction versus human double extraction in evidence synthesis: a randomised controlled trial protocol - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12593456/>
36. Evaluating Language Models for Biomedical Fact-Checking: A Benchmark Dataset for Cancer Variant Interpretation Verification - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12458922/>
37. A General-Purpose VLM Can Teach an Astronomy Foundation Model to Better Recognize Galaxy Morphology - arXiv, <https://arxiv.org/pdf/2608.02300>
38. Galaxy spectroscopy without spectra: Galaxy properties from photometric images with conditional diffusion models - arXiv, <https://arxiv.org/html/2406.18175v1>
39. Identifying lopsidedness in spiral galaxies using a Deep Convolutional Neural Network, <https://arxiv.org/html/2505.19583v2>
40. AI-Powered Detection and Segmentation of Edge-On Galaxies - arXiv, <https://arxiv.org/html/2406.15064v1>
41. A General-Purpose VLM Can Teach an Astronomy Foundation Model to Better Recognize Galaxy Morphology - arXiv, <https://arxiv.org/html/2608.02300v1>
42. Semantic Search for 100M+ Galaxy Images Using AI-Generated Captions - arXiv, <https://arxiv.org/html/2512.11982v2>
43. Inspiring from Galaxies to Green AI in Earth: Benchmarking Energy-Efficient Models for Galaxy Morphology Classification - MDPI, <https://www.mdpi.com/1999-4893/18/7/399>
44. Earth Science Foundation Models: From Perception to Reasoning and Discovery - arXiv, <https://arxiv.org/html/2605.12542v2>
45. State-of-the-Art Review of Artificial Intelligence in Environmental Geophysics and Geotechnical Engineering - Preprints.org, <https://www.preprints.org/manuscript/202605.1511>
46. GitHub - carlos-alberto-silva/satellite-image-deep-learning, <https://github.com/carlos-alberto-silva/satellite-image-deep-learning>
47. Automated Agricultural Land Mapping Using Landsat-8, ArcGIS, and, <https://www.iieta.org/download/file/fid/193729>
48. Full article: A review on the integration of GPR and AI for enhanced road pavement inspection and precise crack detection - Taylor & Francis, <https://www.tandfonline.com/doi/full/10.1080/10298436.2026.2686179>
49. MO-SAM: Testing the reliability and limits of mine feature delineation using Segment Anything Model to democratize mine observation and research — PLOS Sustainability and Transformation, <https://journals.plos.org/sustainabilitytransformation/article?id=10.1371/journal.pstr.0000182>
50. Annotation fatigue: Why human data quality declines over time - Pareto AI, <https://pareto.ai/blog/annotation-fatigue-why-human-data-quality-declines-over-time>
51. Annotation Fatigue in AI: Why Human Data Quality Declines Over Time InnoApps, <https://innoapps.com/blog/annotation-fatigue-in-ai-data-labeling>

52. Top 7 Data Labeling Mistakes That Hurt Your Machine Learning Model Performance, <https://tinkogroup.com/data-labeling-mistakes-ai-performance/>
53. AI-assisted data labeling vs Manual ROI Guide - Tinkogroup, <https://tinkogroup.com/ai-assisted-data-labeling-vs-manual-roi-guide/>
54. Bias in the Loop: How Humans Evaluate AI-Generated Suggestions - arXiv, <https://arxiv.org/html/2509.08514v1>
55. Bias in the Loop: How Humans Evaluate AI-Generated Suggestions - ResearchGate, https://www.researchgate.net/publication/395403018_Bias_in_the_Loop_How_Humans_Evaluate_AI-Generated_Suggestions
56. Automation Bias in Human-AI Collaboration — PDF — Artificial Intelligence - Scribd, <https://www.scribd.com/document/890318722/AutomationBias-Romeo-Conti-25>
57. Trustworthy human computation: a survey - Kyoto University, <https://repository.kulib.kyoto-u.ac.jp/bitstream/2433/294621/1/s10462-024-10974-1.pdf>
58. From Trust to Appropriate Reliance: Measurement Constructs in Human-AI Decision-Making, https://www.researchgate.net/publication/404209597_From_Trust_to_Appropriate_Reliance_Measurement_Constructs_in_Human-AI_Decision-Making
59. Transforming Data Annotation with AI Agents: A Review of Architectures, Reasoning, Applications, and Impact - MDPI, <https://www.mdpi.com/1999-5903/17/8/353>
60. Stuck on Suggestions: Automation Bias, the Anchoring Effect, and the Factors That Shape Them in Computational Pathology - arXiv, <https://arxiv.org/html/2603.11821v2>
61. An Analysis of Active Learning Algorithms using Real-World Crowd-sourced Text Annotations - arXiv, <https://arxiv.org/html/2604.23290>
62. Efficient Data Annotation with Active Learning [Updated] - Labellerr, <https://www.labellerr.com/blog/what-is-active-learning-and-how-does-it-reduce-time-in-data-annotat>
63. Beyond uncertainty in modern active learning for trustworthy AI - Frontiers, <https://www.frontiersin.org/journals/artificial-intelligence/articles/10.3389/frai.2026.1844765/full>
64. Deep Active Learning for Data Mining from Conflict Text Corpora - arXiv, <https://arxiv.org/html/2402.01577v1>
65. Model Uncertainty based Active Learning on Tabular Data using Boosted Trees - arXiv, <https://arxiv.org/pdf/2310.19573>
66. A Novel Active Learning Approach to Label One Million Unknown Malware Variants - arXiv, <https://arxiv.org/html/2507.02959v1>
67. Convergence of Uncertainty Sampling for Active Learning - arXiv, <https://arxiv.org/pdf/2110.15784>
- 2110.15784 Convergence of Uncertainty Sampling for Active Learning - arXiv, <https://arxiv.org/abs/2110.15784>
68. Efficient Human-in-the-Loop Active Learning: A Novel Framework for Data Labeling in AI Systems - arXiv, <https://arxiv.org/pdf/2501.00277>

- 2305.20045 ActiveAED: A Human in the Loop Improves Annotation Error Detection - arXiv, <https://arxiv.org/abs/2305.20045>
69. ActiveAED: A Human in the Loop Improves Annotation Error Detection - ACL Anthology, <https://aclanthology.org/2023.findings-acl.562.pdf>
70. Performance of active learning models for screening prioritization in systematic reviews: a simulation study into the Average Time to Discover relevant records - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC10280866/>
71. Conformal-in-the-Loop for Learning with Imbalanced Noisy Data - arXiv, <https://arxiv.org/html/2411.02281v1>
72. arXiv:2205.12702v3 [cs.CL] 15 Dec 2022, <https://arxiv.org/pdf/2205.12702>
73. Confident Learning-Based Label Correction for Retinal Image Segmentation - MDPI, <https://www.mdpi.com/2075-4418/15/14/1735>
74. Characterizing Label Errors: Confident Learning for Noisy-Labeled Image Segmentation, https://www.researchgate.net/publication/345238962_Characterizing_Label_Errors_Confident_Learning_for_Noisy-Labeled_Image_Segmentation
75. Bias-Aware Mislabeling Detection via Decoupled Confident Learning - arXiv, <https://arxiv.org/html/2507.07216v2>
76. Domain Adaptation for Medical Image Analysis: A Survey - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC9011180/>
77. Handling Out-of-Distribution Data: A Survey - arXiv, <https://arxiv.org/html/2507.21160v1>
78. Out of distribution learning in bioinformatics: advancements and challenges - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12203079/>
- 2410.08000 AHA: Human-Assisted Out-of-Distribution Generalization and Detection - arXiv, <https://arxiv.org/abs/2410.08000>
79. Evaluation of out-of-distribution detection methods for data shifts in single-cell transcriptomics - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12121363/>
- 2306.08328 Distribution Shift Inversion for Out-of-Distribution Prediction - arXiv, <https://arxiv.org/abs/2306.08328>